

# AI基础

## Lecture 2: Problem Solving: Search

Bin Yang

School of Data Science and Engineering

byang@dase.ecnu.edu.cn

[Some slides adapted from Nikita Kitaev, UCB]

# Lecture 1 ILOs

- What is AI?
  - Understand the four approaches
    - Human vs. rationality
    - Behavior vs. thought
  - Beneficial machines
    - Value alignment, AI risks
- AI history
- Agents
  - Key concepts: agent, environment, agent function, agent program, rational agent
  - Performance measure and rational agent
    - Maximize expected performance measure
  - Task environments, and their categorization
    - PEAS, Fully/Partially observable, single-/multi-agent, (non)deterministic, episodic vs. sequential, static vs. dynamic, discrete vs. continuous, and (un)known.
  - Agent programs, and their categorization
    - Simple reflex agents, model-based reflex agents, Goal-based agents, Utility-based agents
    - Learning agents

	Human	Rational
Behavior	<b>Act like people</b> Acting Humanly The Turing test approach	<b>Act rationally</b> The rational agent approach
Thought	<b>Think like people</b> Thinking humanly The cognitive modeling approach	<b>Think rationally</b> The “laws of thought” approach




Figure 2.1 Agents interact with environments through sensors and actuators.

# Lecture 2 ILOs

- Problem-solving agent
  - What is a problem-solving agent
  - Search problems and solutions
- Example problems
  - Standardized problems
  - Real-world problems
- Search algorithms
  - Uninformed search vs. informed search
  - Performance measure: completeness, optimality, time complexity, space complexity
  - Best-First-Search
- Uninformed search
  - Breath-First-Search
  - Uniform cost search/Dijkstra's algorithm
  - Depth-First-Search
  - Iterative deepening search
  - Bidirectional search

# Outline

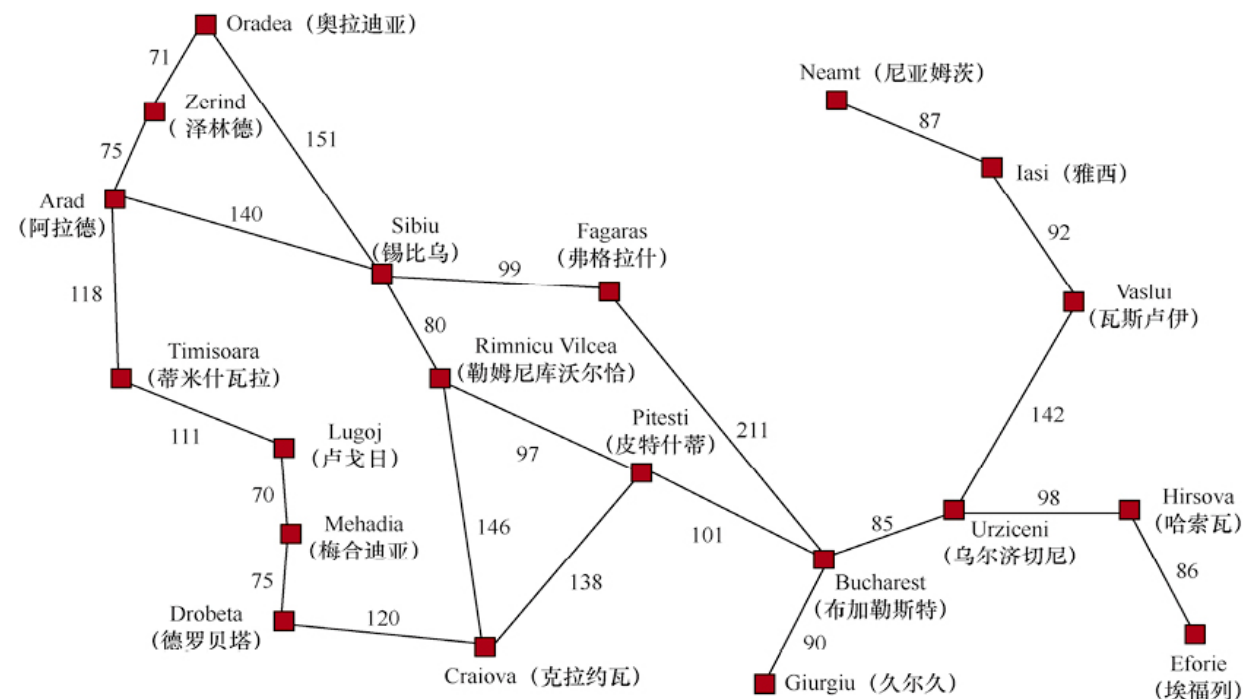
- Problem-solving agent
- Search problem and solutions
- Uninformed Search Algorithms

# Solving problems by searching

- When it is useful?
  - When the correct action to take is not immediately obvious
- Problem-solving agent
  - An agent may need to plan ahead: a sequence of actions that form a path to a goal state
- Search
  - Computational process that a problem-solving agent takes
  - Informed searching vs. uninformed searching
    - Whether the agent can estimate how far it is from the goal
- Environment, simplest
  - Episodic, single agent, fully observable, deterministic, static, discrete, and known.
  - We will also show a motivating example on more difficult environments.

# An example

- Arad to Bucharest
  - Known: with map
    - Three roads leaving Arad
  - Unknown: no map



# Problem-solving agent

- Goal formulation
  - Bucharest.
  - Goals organize behavior by limiting the objectives and hence the actions to be considered.
- Problem formulation
  - States and actions necessary to reach the goal
- Search
  - The agent simulates sequences of actions in its model, searching until it finds a sequence of actions that reaches the goal. This is called a solution.
- Execution
  - Execute the actions in the solution

# Open-loop and closed-loop system

- Fully observable, deterministic, and known environment
  - The solution to any problem is a **fixed sequence of actions**
  - **Open-loop**, ignore percepts while executing breaks the loop between agent and environment.
  - Example: shortest path
- Partially observable, non-deterministic
  - Actions depends on what percepts arrive
  - Example: traffic based fastest path
  - The solution is a **strategy**
    - Different future actions based on percepts
  - **Closed-loop**

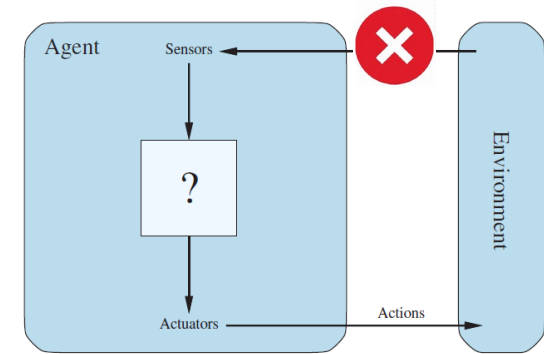
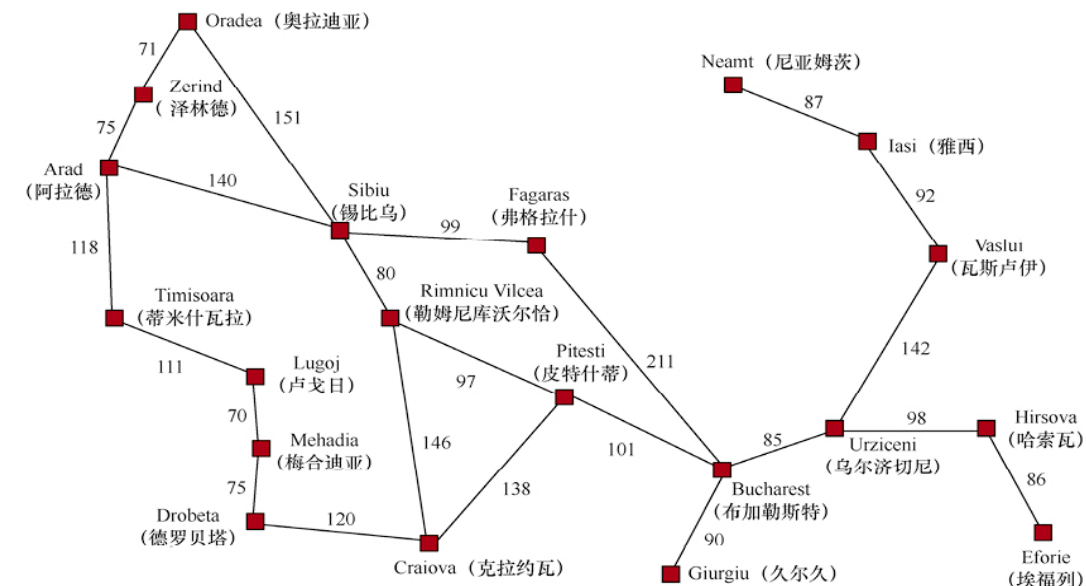


Figure 2.1 Agents interact with environments through sensors and actuators.

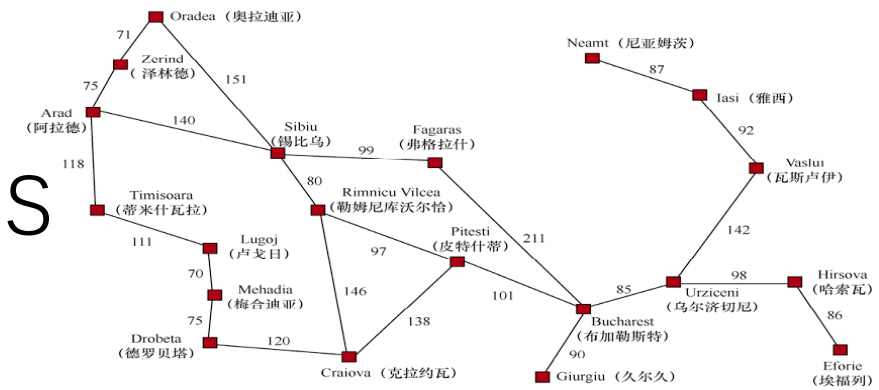




# Outline


- Problem-solving agent
- Search problem and solutions
- Uninformed Search Algorithms

# Search problems and solutions



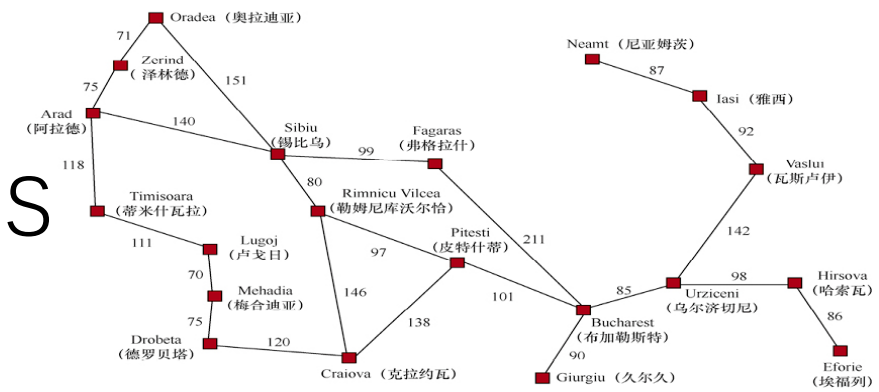
- State space
  - A set of possible states that the environment can be in.
- Initial state
  - Where the agent starts.
  - E.g., Arad
- Goal states
  - One goal vs. many goals
  - E.g., Bucharest vs. no dirt in any location (vacuum-cleaner)
  - Is-Goal method

# Search problems and solutions



- Actions
  - The actions available to the agent at state  $s$ .
  - $\text{Actions}(s)$  returns a finite set of actions that can be executed in state  $s$ .
  - $\text{Action}(\text{Arad}) = \{\text{ToSibiu}, \text{ToTimisoara}, \text{ToZerind}\}$
- Transition model
  - $\text{Result}(s, a)$ : returns the state from doing action  $a$  in state  $s$ .
  - $\text{Result}(\text{Arad}, \text{ToZerind}) = \text{Zerind}$
- Action cost function
  - $\text{Action-Cost}(s, a, s')$ : numeric cost of applying action  $a$  in state  $s$  to reach a new state  $s'$
  - $\text{Action-Cost}(\text{Arad}, \text{ToZerind}, \text{Zerind}) = 75$

# Search problems and solutions



- The state space can be modeled as a graph
  - **State-space-graph**
  - Nodes: states
  - Edges: actions
  - Cities, nodes, states
  - Roads, edges, actions (one road corresponds to two actions, one in each direction)
- Path: a sequence of actions
- Solution: a path from the initial state to a goal state
- Optimal solution: the solution with the lowest path cost
  - Assume action costs are additive. The cost of a path is the sum of the costs of the individual actions in the path.
  - This assumption may not be always true.


# Example Problems

- Standardized problems vs. real-world problems
  - Concise, abstract, exact description. Suitable for benchmarking performance of different problem-solving algorithms
  - Whose solution people actually can use, formulation is idiosyncratic
- Standardized problems
  - Grid world, vacuum world
- Real-world problems
  - Routing-finding
  - Touring problems
  - VLSI layout
  - Robot navigation

# Grid world



- 2D rectangular array of cells
- An agent can move from cell to cell
- Cells contain objects
- The agent can operate on the objects

- Vacuum world




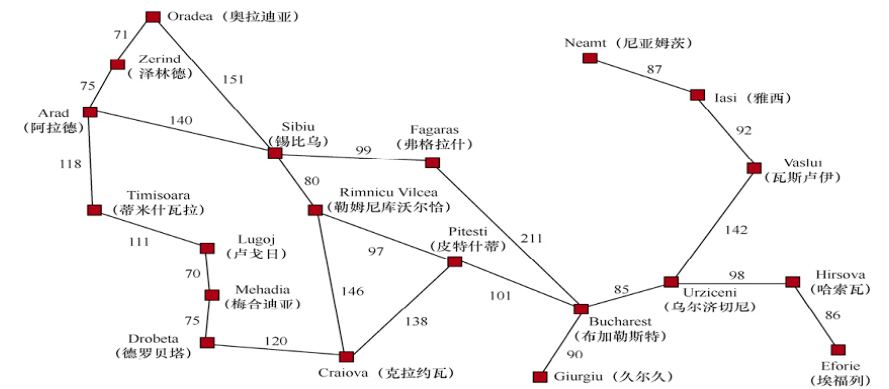
# Vacuum world

- States:
  - Which objects (agent/dirt) are in which cells
  - The agent can be in one of the two cells
  - Each cell can either contain dirt or not
  - Quiz: how many states do we have?
  - 2 cells example: 8 states
  - $n$  cells:  $n \cdot 2^n$  states
- Initial state
  - Any state in the 8 states
- Actions
  - Suck, left, right
- Transition model
- Goal states: every cell is clean
- Action cost: constant
- State-space-graph
  - We can rarely build full graph in memory (it's too big) when  $n$  is big



# Search algorithms


- Input: a search problem
- Output: a solution, or an indication of failure (no solution exists)
- Algorithms that superimpose a **search tree** over the state-space graph
  - Nodes: states
  - Edges: actions
  - Root: initial state





# State Space Graphs vs. Search Trees


## State Space Graph



*Each NODE in the search tree is an entire PATH in the state space graph.*


*We construct both on demand – and we construct as little as possible.*

## Search Tree



# Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?




# Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:




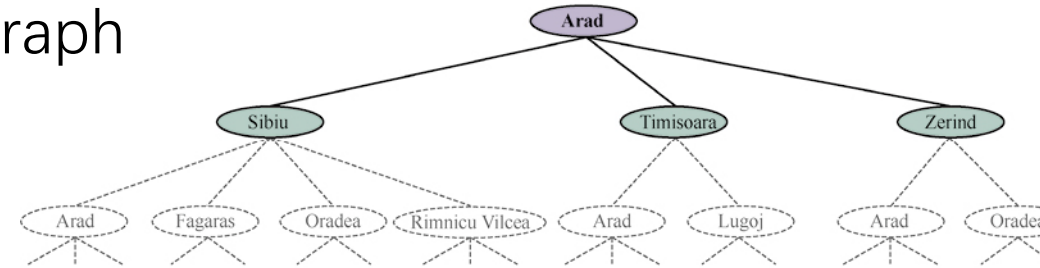
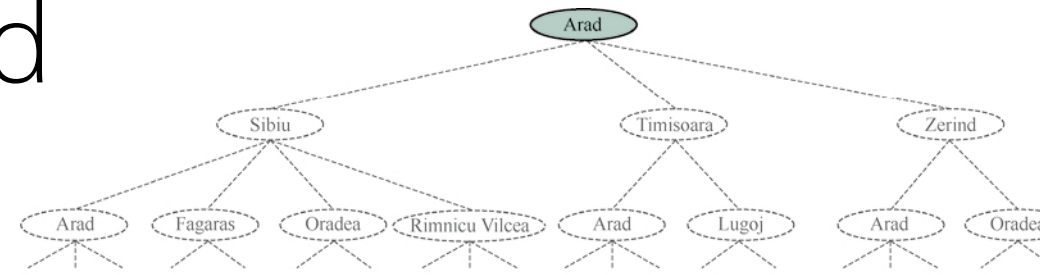



How big is its search tree (from S)?



Important: Lots of repeated structure in the search tree!

# Frontier, reached, expanded

- Frontier: reached, but not yet expanded
  - Separates two regions of the state-space graph
    - Interior region: expanded
    - Exterior region: not yet reached



# Best-first search

- Which node from the frontier to expand next?
  - We choose a node  $n$  with the minimum value of evaluation function  $f(n)$
- Each iteration
  - Choose a node on the frontier with minimum  $f(n)$  to expand/stop
    - Priority queue
  - The node has not been reached before, or re-added

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure  
  node  $\leftarrow$  NODE(STATE=problem.INITIAL)  
  frontier  $\leftarrow$  a priority queue ordered by f, with node as an element  
  reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node  
  while not IS-EMPTY(frontier) do  
    node  $\leftarrow$  POP(frontier)  
    if problem.IS-GOAL(node.STATE) then return node  
    for each child in EXPAND(problem, node) do  
      s  $\leftarrow$  child.STATE  
      if s is not in reached or child.PATH-COST < reached[s].PATH-COST then  
        reached[s]  $\leftarrow$  child  
        add child to frontier  
  return failure
```

```
function EXPAND(problem, node) yields nodes  
  s  $\leftarrow$  node.STATE  
  for each action in problem.ACTIONS(s) do  
    s'  $\leftarrow$  problem.RESULT(s, action)  
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')  
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

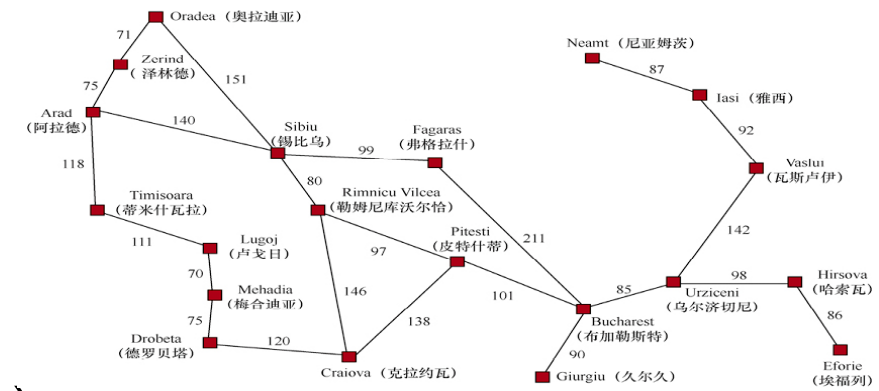
# Search data structures

- Node
  - State
  - Parent: the node in the tree that generated this node
  - Action: the action applied on the parent node
  - Path-cost: total cost of path from the initial state to this node
- Frontier, maintain in a “queue”
  - Priority queue, best-first search
  - FIFO queue, breadth-first search
  - LIFO queue (stack), depth-first search

# Measuring problem-solving performance

- Completeness
  - The algorithm is guaranteed to find a solution when there is one, and to correctly report failure when there is not.
- Cost optimality
  - Find a solution with the lowest path cost among all solutions.
- Time complexity
  - How long does it take to find a solution. In terms of the number of states and actions.
- Space complexity
  - How much memory is needed.
- Complexity
  - When the graph is an explicit data structure
    - Map: V, E
  - Implicit graph: transition model, actions
    - d, depth, or the number of actions in an optimal solution;
    - m, the maximum number of actions in any path;
    - b, branching factor, number of successors of a node that need to be considered


# Uninformed search strategies



- No idea on how close a state is to the goal(s).
- Consider the example from Arad to Bucharest
- Uninformed agent
  - No knowledge of Romanian geography, Zerind or Sibiu?
- Informed agent
  - Knows the location of each city and knows Sibiu is closer to the goal.




# Breadth-First Search (BFS)



# Breadth-first search (BFS)

- When all actions have the same cost
- Root node is expanded
- Then, all the successors of the root node are expanded
- Then, the successors of the successors are expanded





- When using the Best-First-Search,  $f(n)$  is the depth of the node
- More efficient implementation of BFS is to use a FIFO queue.

# Breadth-First Search

*Strategy: expand a shallowest node first*

*Implementation: Frontier is a FIFO queue*



# Search Algorithm Properties

- Complete: Guaranteed to find a solution if one exists
- Optimal: Guaranteed to find the least cost path
- Time complexity?
- Space complexity?

- Cartoon of search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^m)$

## Geometric series

For real  $x \neq 1$ , the summation


$$\sum_{k=0}^n x^k = 1 + x + x^2 + \dots + x^n$$

is a *geometric* or *exponential series* and has the value

$$\sum_{k=0}^n x^k = \frac{x^{n+1} - 1}{x - 1} . \quad (\text{A.5})$$


When the summation is infinite and  $|x| < 1$ , we have the infinite decreasing geometric series

$$\sum_{k=0}^{\infty} x^k = \frac{1}{1 - x} . \quad (\text{A.6})$$



# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
  - Processes all nodes above d shallowest solution
  - Let depth of the shallowest solution be  $d$
  - Search takes time  $O(b^d)$
- How much space does the frontier take?
  - Has roughly the last tier, so  $O(b^d)$
- Is it complete?
  - $d$  must be finite if a solution exists, so yes!
- Is it optimal?
  - Only if costs of all actions are the same, e.g., 1





# Dijkstra's algorithm or uniform-cost search

- When actions have different costs
- Best-first-search:  $f(n)$  is the cost of path from the root to the current node  $n$ .
- Breadth first search
  - First depth 1, second depth 2, third depth 3, ...
  - When exploring nodes in depth  $d$ , all nodes with depth  $< d$  have been explored.
- Uniform cost search
  - When exploring nodes with cost  $c$ , all nodes with cost  $< c$  have been explored.

# Uniform Cost Search


*Strategy: expand a  
cheapest node first*

*Frontier is a priority queue  
(priority: cumulative cost)*



# Uniform Cost Search (UCS) Properties

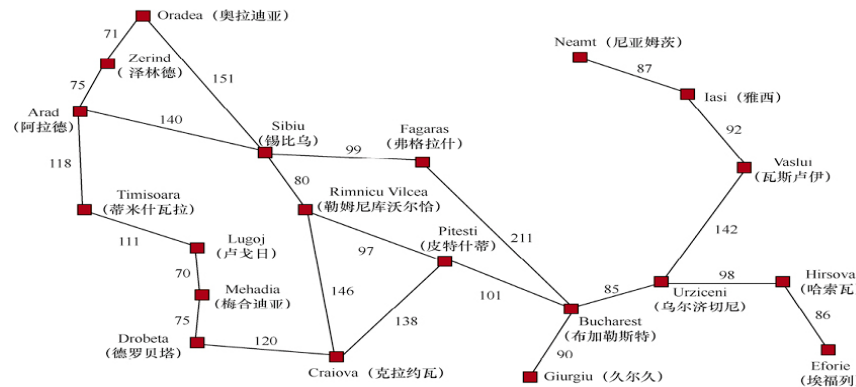


- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and action cost at least  $\epsilon$ , then the “effective depth” is roughly  $C^*/\epsilon$
  - Takes time  $O(b^{C^*/\epsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\epsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum action cost is positive, yes!
- Is it optimal?
  - Yes!
  - The first solution it finds will have a cost that is at least as low as the cost of any other node in frontier.







# Uniform Cost Issues

- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We'll fix that soon!





# Bidirectional search


- Search forward from the initial state and search backwards from the goal state(s).
- When the two frontiers collide, we have a solution.
- Grow a ball around start and goal when goal is scanned
- Grow a ball around each end (start and goal) until they “meet”.
- Can be used with other speed-up methods, e.g., contraction hierarchies.



# Bidirectional search



# Uncertainty, episodic vs. sequential, deterministic vs. nondeterministic



- Example:  $P = \langle e_5, e_6 \rangle$ 
  - Trajectory 1: 10 s, 20 s
  - Trajectory 2: 15 s, 25 s

$e_5$	Cost	Prob
	10	0.5
	15	0.5

$e_6$	Cost	Prob
	20	0.5
	25	0.5

- The distribution of a path is computed by summing the distributions of edges while assuming they are independent.

Independence  
 $\langle e_5, e_6 \rangle$

Cost	Prob
10, 20	0.25
15, 20	0.25
10, 25	0.25
15, 25	0.25

$P$	
Cost	Prob
30	0.25
35	0.50
40	0.25

Dependence

$\langle e_5, e_6 \rangle$	
Cost	Prob
10, 20	0.5
15, 25	0.5

$P$	
Cost	Prob
30	0.5
40	0.5

Path-centric weight

Joint distribution

Yang, Dai, Guo, Jensen, Hu. PACE: A Path-Centric Paradigm For Stochastic Path Finding. The VLDB Journal 27(2): 153-178 (2018).

Dai, Yang, Guo, Jensen, Hu. Path Cost Distribution Estimation Using Trajectory Data. PVLDB 10(3): 85-96 (2016).

Yang, Guo, Jensen, Kaul, Shang. Stochastic Skyline Route Planning Under Time-Varying Uncertainty. ICDE 2014, 136-147, 2014.

# Depth-First Search



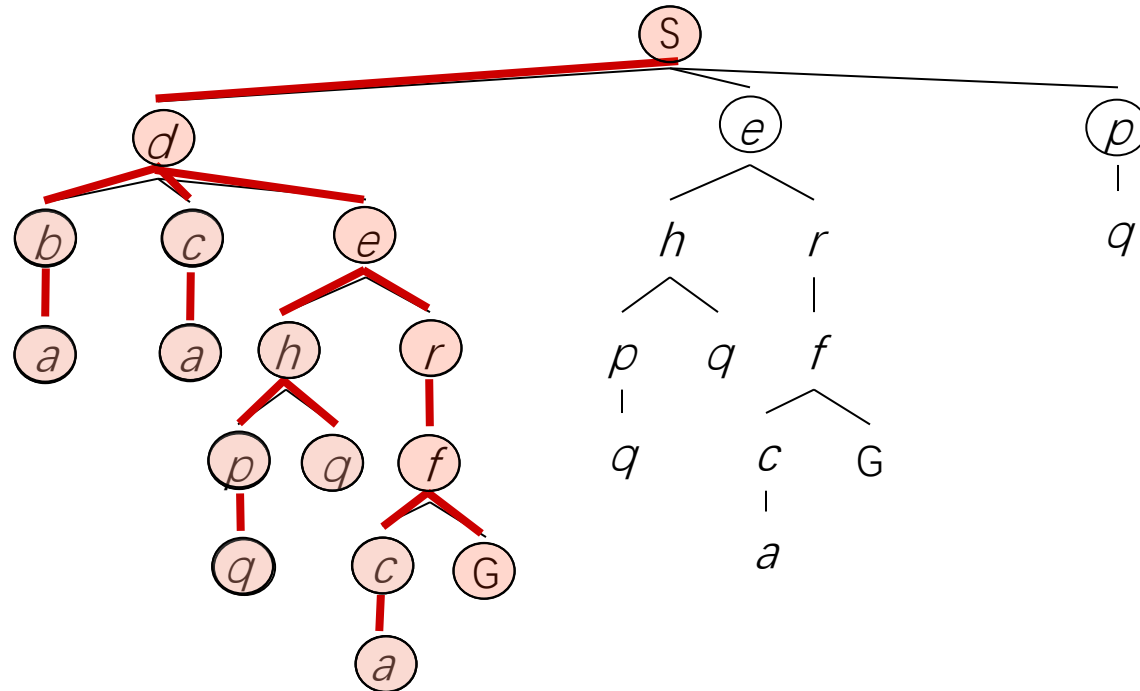

# Depth-first search (DFS)

- When using the Best-First-Search,  $f(n)$  is the negative of the depth of the node
- Search proceeds immediately to the deepest level of the search tree, where the nodes have no successors.
  - Then, backtrack to the next deepest node that still has unexpanded successors.
- Not cost-optimal, it returns the first solution it finds.

# Depth-First Search


*Strategy: expand a deepest node first*

*Implementation:  
Frontier is a LIFO stack*



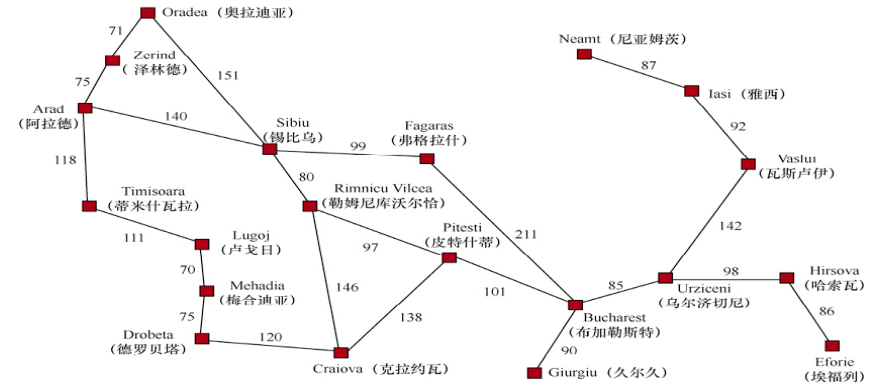
# Depth-First Search (DFS) Properties

- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the frontier take?
  - Only has siblings on path to root, so  $O(bm)$
  - Good space complexity
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost
- Not complete and not optimal, then why we care DFS?







# Depth-limited DFS



- Given a depth limit  $l$ , and treat all nodes at depth  $l$  as if they had no successors.
- Time complexity  $O(b^l)$
- Space complexity  $O(bl)$
- Romania map example
  - 20 cities in total, so  $l=19$  is a valid limit.
  - One city can be reached from any other city in at most 9 actions.  $l=9$ .
  - Diameter of the state-space graph

# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. ...



# Iterative Deepening


- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages

- Run a DFS with depth limit 1. If no solution...
- Run a DFS with depth limit 2. If no solution...
- Run a DFS with depth limit 3. ...

- Isn't that wastefully redundant?

- Generally most work happens in the lowest level searched, so not so bad!
- Time complexity : like BFS, still  $O(b^d)$
- $b=10, d=5$
- BFS:  $10+10^2+10^3+10^4+10^5=111,110$
- IDS:  $5*10+4*10^2+3*10^3+2*10^4+10^5=123,450$

- Space complexity, like DFS
- Optimal, if all actions have the same cost, like BFS
- Complete, finite acyclic state spaces



# Uninformed Search Summary

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No	No	Yes <sup>1</sup>	Yes <sup>1,4</sup>
Optimal cost?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>	Yes <sup>3,4</sup>
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$

**Figure 3.15** Evaluation of search algorithms.  $b$  is the branching factor;  $m$  is the maximum depth of the search tree;  $d$  is the depth of the shallowest solution, or is  $m$  when there is no solution;  $\ell$  is the depth limit. Superscript caveats are as follows: <sup>1</sup> complete if  $b$  is finite, and the state space either has a solution or is finite. <sup>2</sup> complete if all action costs are  $\geq \epsilon > 0$ ; <sup>3</sup> cost-optimal if action costs are all identical; <sup>4</sup> if both directions are breadth-first or uniform-cost.

# Lecture 2 ILOs

- Problem-solving agent
  - What is a problem-solving agent
  - Search problems and solutions
    - State space, initial state, goal states, action, transition model, action cost function
    - A solution: a path from the initial state to a goal state
- Example problems
  - Standardized problems
  - Real-world problems
- Search algorithms
  - Uninformed search vs. informed search
  - Performance measure: completeness, optimality, time complexity, space complexity
  - Best-First-Search: Which node from the frontier to expand next? Node with minimum  $f(n)$ .
- Uninformed search
  - Breath-First-Search
  - Uniform cost search/Dijkstra's algorithm
  - Bidirectional search
  - Depth-First-Search
  - Iterative deepening search